

```
mkoken@linux60813:~  
[mkoken@linux60813 ~]$ hadoop jar $HADOOP_PREFIX/share/hadoop/mapreduce/hadoop-mapreduce-examples*.jar wordcount war-and-peace.txt wordcount-out  
16/05/31 11:56:39 INFO client.RMProxy: Connecting to ResourceManager at name1.hadoop.dc.engr.scu.edu/10.128.0.201:8032  
16/05/31 11:56:39 INFO input.FileInputFormat: Total input paths to process : 1  
16/05/31 11:56:39 INFO mapreduce.JobSubmitter: number of splits:1  
16/05/31 11:56:40 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1458785012989_1730  
16/05/31 11:56:40 INFO impl.YarnClientImpl: Submitted application application_1458785012989_1730  
16/05/31 11:56:40 INFO mapreduce.Job: The url to track the job: http://name1.hadoop.dc.engr.scu.edu:8088/proxy/application_1458785012989_1730/  
16/05/31 11:56:40 INFO mapreduce.Job: Running job: job_1458785012989_1730  
16/05/31 11:56:44 INFO mapreduce.Job: Job job_1458785012989_1730 running in uber mode : false  
16/05/31 11:56:44 INFO mapreduce.Job:  map 0% reduce 0%  
16/05/31 11:56:52 INFO mapreduce.Job:  map 100% reduce 0%  
16/05/31 11:56:57 INFO mapreduce.Job:  map 100% reduce 26%  
16/05/31 11:56:58 INFO mapreduce.Job:  map 100% reduce 88%  
16/05/31 11:56:59 INFO mapreduce.Job:  map 100% reduce 100%  
16/05/31 11:56:59 INFO mapreduce.Job: Job job_1458785012989_1730 completed successfully  
16/05/31 11:56:59 INFO mapreduce.Job: Counters: 49  
    File System Counters  
        FILE: Number of bytes read=465077  
        FILE: Number of bytes written=12156740  
        FILE: Number of read operations=0  
        FILE: Number of large read operations=0  
        FILE: Number of write operations=0  
        HDFS: Number of bytes read=3226769  
        HDFS: Number of bytes written=467832  
        HDFS: Number of read operations=291  
        HDFS: Number of large read operations=0  
        HDFS: Number of write operations=192  
    Job Counters  
        Launched map tasks=1  
        Launched reduce tasks=96  
        Rack-local map tasks=1  
        Total time spent by all maps in occupied slots (ms)=2756  
        Total time spent by all reduces in occupied slots (ms)=348008  
        Total time spent by all map tasks (ms)=2756  
        Total time spent by all reduce tasks (ms)=348008  
        Total vcore-seconds taken by all map tasks=2756  
        Total vcore-seconds taken by all reduce tasks=348008  
        Total megabyte-seconds taken by all map tasks=2822144  
[0] 0:[tmux]* 1:bash- "mkoken@linux60813:~" 11:58 31-May-16
```

Launched map tasks=1
Launched reduce tasks=96
Rack-local map tasks=1
Total time spent by all maps in occupied slots (ms)=2756
Total time spent by all reduces in occupied slots (ms)=348008
Total time spent by all map tasks (ms)=2756
Total time spent by all reduce tasks (ms)=348008
Total vcore-seconds taken by all map tasks=2756
Total vcore-seconds taken by all reduce tasks=348008
Total megabyte-seconds taken by all map tasks=2822144
Total megabyte-seconds taken by all reduce tasks=356360192

Map-Reduce Framework

Map input records=65007
Map output records=566316
Map output bytes=5478297
Map output materialized bytes=464693
Input split bytes=135
Combine input records=566316
Combine output records=42065
Reduce input groups=42065
Reduce shuffle bytes=464693
Reduce input records=42065
Reduce output records=42065
Spilled Records=84130
Shuffled Maps =96
Failed Shuffles=0
Merged Map outputs=96
GC time elapsed (ms)=2237
CPU time spent (ms)=91730
Physical memory (bytes) snapshot=25875963904
Virtual memory (bytes) snapshot=155323228160
Total committed heap usage (bytes)=48984227840

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=3226634

File Output Format Counters

Bytes Written=467832

[mkoken@linux60813 ~]\$

mkoken@linux60813:~

```
"Adorable!      1
"Again!"        1
"Aye,   1
"Before 2
"Do,   1
"Everything's   2
"Hold   1
"Non,   2
"Posterity     1
"Saber  1
"Uncle" 35
"Wostov!      1
"a      8
"all     4
"but    68
"plenary"     1
'One     2
(does    1
809      1
A       373
Across   1
Ahead    1
Alexander--though 1
All      223
Anisya   14
Antonov." 1
Auerstadt. 1
Bazdeev, 2
Before   51
Belliard. 1
Bennigsen: 2
But      829
Catiche,  5
Could    5
Daniel   27
Directory, 2
Dron.    7
Duke,    3
Elizabeth? 1
F3.      1
Ferdinand  6
Fortune  1
French--having 1
French--in  1
```

:

[0] 0: bash* 1: bash-

"mkoken@linux60813:~" 12:00 31-May-16

mkoken@linux60813:~

16/05/31 12:06:58 INFO impl.YarnClientImpl: Submitted application application_1458785012989_1732

16/05/31 12:06:58 INFO mapreduce.Job: The url to track the job: http://name1.hadoop.dc.engr.scu.edu:8088/proxy/application_1458785012989_1732/

16/05/31 12:06:58 INFO mapreduce.Job: Running job: job_1458785012989_1732

16/05/31 12:07:02 INFO mapreduce.Job: Job job_1458785012989_1732 running in uber mode : false

16/05/31 12:07:02 INFO mapreduce.Job: map 0% reduce 0%

16/05/31 12:07:12 INFO mapreduce.Job: map 25% reduce 0%

16/05/31 12:07:13 INFO mapreduce.Job: map 61% reduce 0%

16/05/31 12:07:14 INFO mapreduce.Job: map 100% reduce 0%

16/05/31 12:07:24 INFO mapreduce.Job: map 100% reduce 82%

16/05/31 12:07:25 INFO mapreduce.Job: map 100% reduce 97%

16/05/31 12:07:26 INFO mapreduce.Job: map 100% reduce 100%

16/05/31 12:07:26 INFO mapreduce.Job: Job job_1458785012989_1732 completed successfully

16/05/31 12:07:26 INFO mapreduce.Job: Counters: 50

File System Counters

FILE: Number of bytes read=4470698469

FILE: Number of bytes written=8886852534

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=10000013248

HDFS: Number of bytes written=10000000000

HDFS: Number of read operations=360

HDFS: Number of large read operations=0

HDFS: Number of write operations=48

Job Counters

Launched map tasks=96

Launched reduce tasks=24

Data-local map tasks=89

Rack-local map tasks=7

Total time spent by all maps in occupied slots (ms)=542919

Total time spent by all reduces in occupied slots (ms)=212839

Total time spent by all map tasks (ms)=542919

Total time spent by all reduce tasks (ms)=212839

Total vcore-seconds taken by all map tasks=542919

Total vcore-seconds taken by all reduce tasks=212839

Total megabyte-seconds taken by all map tasks=555949056

Total megabyte-seconds taken by all reduce tasks=217947136

Map-Reduce Framework

Map input records=100000000

Map output records=100000000

Map output bytes=10200000000

[0] 0:[tmux]* 1: bash-

"mkoken@linux60813:~" 12:08 31-May-16

mkoken@linux60813:~

[0/357]

Rack-local map tasks=7
Total time spent by all maps in occupied slots (ms)=542919
Total time spent by all reduces in occupied slots (ms)=212839
Total time spent by all map tasks (ms)=542919
Total time spent by all reduce tasks (ms)=212839
Total vcore-seconds taken by all map tasks=542919
Total vcore-seconds taken by all reduce tasks=212839
Total megabyte-seconds taken by all map tasks=555949056
Total megabyte-seconds taken by all reduce tasks=217947136

Map-Reduce Framework

Map input records=100000000
Map output records=100000000
Map output bytes=10200000000
Map output materialized bytes=4402111541
Input split bytes=13248
Combine input records=0
Combine output records=0
Reduce input groups=100000000
Reduce shuffle bytes=4402111541
Reduce input records=100000000
Reduce output records=100000000
Spilled Records=200000000
Shuffled Maps =2304
Failed Shuffles=0
Merged Map outputs=2304
GC time elapsed (ms)=9007
CPU time spent (ms)=634600
Physical memory (bytes) snapshot=70237343744
Virtual memory (bytes) snapshot=191057907712
Total committed heap usage (bytes)=68918706176

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=100000000000

File Output Format Counters

Bytes Written=100000000000

16/05/31 12:07:26 INFO terasort.TeraSort: done

[mkoken@linux60813 ~]\$

[0] 0:[tmux]* 1: bash-

"mkoken@linux60813:~" 12:10 31-May-16

```
[mkoken@linux60813 ~]$ hadoop fs -count -q .
      50000      49842      483183820800      453172734872      9      1
49      20003694719 .
[mkoken@linux60813 ~]$
```

[0] 0: bash* 1: bash- "mkoken@linux60813:~" 12:11 31-May-16

```
[mkoken@linux60813 ~]$ hadoop fs -du -s -h .
18.6 G 28.0 G .
[mkoken@linux60813 ~]$
```

[0] 0: bash* 1: bash- "mkoken@linux60813:~" 12:10 31-May-16